

Kubernetes Logging

Table of Contents

What are the different types of logs in a Kubernetes cluster, and what security-relevant information does each provide?	1
What are Kubernetes audit logs, how do you configure them, and what audit policy should you implement for security monitoring?	6
How do you integrate Kubernetes logs (including control plane logs) with Microsoft Sentinel without using Grafana?	14
How do you integrate Kubernetes logs with Microsoft Sentinel using Grafana Loki as an intermediary, and what are the advantages of this approach?	24
What are the key security events you should detect and alert on from Kubernetes control plane logs in Sentinel?	35

What are the different types of logs in a Kubernetes cluster, and what security-relevant information does each provide?

Kubernetes generates multiple log streams providing comprehensive visibility into cluster operations, security events, and application behavior. Understanding these log types is essential for effective security monitoring.

Kubernetes log categories:

1. Control Plane Logs (most security-critical):

kube-apiserver logs:

- **What it logs** - all API requests to the cluster, authentication and authorization events, admission controller decisions, resource creation/modification/deletion, API access patterns.
- **Security value** - who accessed what resources and when, unauthorized access attempts, privilege escalation attempts, suspicious API usage patterns, compliance audit trail.

Example log entry:

```
{  
  "kind": "Event",  
  "apiVersion": "audit.k8s.io/v1",  
  "level": "Metadata",  
  "auditID": "a1b2c3d4-e5f6-7890-abcd-ef1234567890",  
  "stage": "ResponseComplete",  
  "user": {  
    "name": "root",  
    "uid": "0",  
    "groups": []  
  },  
  "object": {  
    "kind": "Pod",  
    "apiVersion": "v1",  
    "resourceVersion": "1234567890",  
    "name": "nginx",  
    "namespace": "default",  
    "uid": "1234567890",  
    "resource": "pods",  
    "selfLink": "/api/v1/namespaces/default/pods/nginx"  
  },  
  "action": "CREATE",  
  "verb": "create",  
  "target": {  
    "kind": "Pod",  
    "apiVersion": "v1",  
    "resourceVersion": "1234567890",  
    "name": "nginx",  
    "namespace": "default",  
    "uid": "1234567890",  
    "resource": "pods",  
    "selfLink": "/api/v1/namespaces/default/pods/nginx"  
  },  
  "related": {  
    "kind": "Event",  
    "apiVersion": "v1",  
    "resourceVersion": "1234567890",  
    "name": "nginx",  
    "namespace": "default",  
    "uid": "1234567890",  
    "resource": "events",  
    "selfLink": "/api/v1/namespaces/default/events/nginx"  
  },  
  "log": "2023-01-01T12:00:00Z [INFO] Creating pod nginx in namespace default",  
  "source": "kube-apiserver",  
  "component": "apiserver",  
  "host": "192.168.1.1",  
  "version": "1.23.5",  
  "severity": "Info",  
  "logger": "audit",  
  "ts": "2023-01-01T12:00:00Z",  
  "file": "audit.log",  
  "line": 1234567890}
```

```

    "requestURI": "/api/v1/namespaces/production/secrets",
    "verb": "list",
    "user": {
        "username": "alice@example.com",
        "groups": ["developers", "system:authenticated"]
    },
    "sourceIPs": ["203.0.113.45"],
    "userAgent": "kubectl/v1.28.0",
    "objectRef": {
        "resource": "secrets",
        "namespace": "production",
        "apiVersion": "v1"
    },
    "responseStatus": {
        "code": 200
    },
    "requestReceivedTimestamp": "2024-01-20T10:30:00.123456Z",
    "stageTimestamp": "2024-01-20T10:30:00.234567Z"
}

```

Security events to monitor:

- Secret access (ConfigMaps, Secrets)
- Role/RoleBinding modifications
- Exec into pods
- Port-forward connections
- Privileged pod creation
- ServiceAccount token creation
- Admission webhook denials

kube-controller-manager logs:

- **What it logs** - controller operations (ReplicaSet, Deployment reconciliation), garbage collection events, node lifecycle events, persistent volume operations.
- **Security value** - unauthorized controller operations, resource quota violations, suspicious resource creation patterns, node compromise indicators.

kube-scheduler logs:

- **What it logs** - pod scheduling decisions, node affinity/anti-affinity, resource allocation, scheduling failures.
- **Security value** - unusual scheduling patterns, attempts to schedule on specific nodes, resource exhaustion attacks.

etcd logs:

- **What it logs** - distributed key-value store operations, cluster state changes, backup/restore

operations.

- **Security value** - direct etcd access (should be none except kube-apiserver), data corruption attempts, unauthorized state modifications.

cloud-controller-manager logs (cloud-specific):

- **What it logs** - cloud provider interactions, load balancer provisioning, node lifecycle in cloud, persistent volume provisioning.
- **Security value** - unauthorized cloud resource creation, suspicious load balancer configurations.

2. Node-level logs:

kubelet logs:

- **What it logs** - pod lifecycle on the node (starting, stopping), image pulls, container runtime interactions, volume mounts, health checks.
- **Security value** - privileged container creation, hostPath volume usage, suspicious image pulls, failed pod starts, resource exhaustion on node.

Example kubelet log:

```
I0120 10:30:00.123456 kubelet.go:1234] Creating pod: production/webapp-abc123
W0120 10:30:01.234567 kubelet.go:5678] Pod production/webapp-abc123 attempted to mount
hostPath /etc, blocked by admission
E0120 10:30:02.345678 kubelet.go:9012] Failed to pull image
"malicious.registry.com/backdoor:latest": unauthorized
```

kube-proxy logs:

- **What it logs** - network proxy operations, service endpoint updates, iptables/IPVS rule changes.
- **Security value** - network policy bypasses, suspicious service access, port scanning detection.

Container runtime logs (containerd, Docker, CRI-O):

- **What it logs** - container lifecycle events, image operations, runtime errors.
- **Security value** - container escape attempts, runtime vulnerabilities exploited, image integrity violations.

3. Application logs:

Pod/Container logs (stdout/stderr):

- **What it logs** - application-specific logging, business logic events, errors and exceptions.
- **Security value** - application-level attacks (SQL injection, XSS), authentication failures, suspicious user behavior, data access patterns.

Sidecar logs:

- **What it logs** - service mesh traffic (Istio, Linkerd), logging agents (Fluentd, Fluent Bit), security scanning (Falco).

4. Audit logs (critical for security):

Kubernetes Audit Logs:

- **What they capture** - comprehensive audit trail of all API server interactions, configurable levels (None, Metadata, Request, RequestResponse), policy-driven (what to log).
- **Security value** - complete forensic record, compliance requirements (PCI-DSS, HIPAA, SOC 2), incident investigation, insider threat detection.

Audit policy levels:

```

apiVersion: audit.k8s.io/v1
kind: Policy
rules:
  # Log all secret access at RequestResponse level (full payload)
  - level: RequestResponse
    resources:
      - group: ""
        resources: ["secrets"]

  # Log all authentication/authorization at Metadata level
  - level: Metadata
    omitStages:
      - RequestReceived
    verbs: ["create", "update", "patch", "delete"]

  # Don't log read-only requests to non-sensitive resources
  - level: None
    verbs: ["get", "list", "watch"]
    resources:
      - group: ""
        resources: ["configmaps", "endpoints"]

  # Log everything else at Request level (no response body)
  - level: Request

```

5. Cluster add-on logs:

DNS logs (CoreDNS):

- **What they log** - DNS queries and responses, cache behavior, query failures.
- **Security value** - DNS tunneling detection, C2 communication, data exfiltration via DNS, malicious domain access.

Ingress controller logs:

- **What they log** - HTTP/HTTPS requests, TLS handshakes, routing decisions, rate limiting.
- **Security value** - web attacks (OWASP Top 10), DDoS attempts, suspicious user agents, geographic anomalies.

Network policy logs:

- **What they log** - allowed/denied connections, policy violations.
- **Security value** - lateral movement attempts, unauthorized service access, network reconnaissance.

Log security priorities:

Highest priority (must monitor):

1. kube-apiserver audit logs (all API access)
2. Secret/ConfigMap access
3. Role/RoleBinding changes
4. Privileged pod creation
5. Exec/port-forward sessions

High priority:

1. Authentication failures
2. Admission webhook denials
3. Image pull failures
4. Suspicious network connections
5. Node kubelet errors

Medium priority:

1. Application errors
2. DNS queries
3. Ingress logs
4. Resource quota violations

Storage considerations:

- Control plane logs: 1-5 GB/day per cluster
- Node logs: 500 MB - 2 GB/day per node
- Application logs: Varies widely (1-100 GB/day)
- Audit logs with RequestResponse: 10-50 GB/day (high verbosity)

Retention recommendations:

- Hot storage (immediate query): 30-90 days

- Warm storage (archive query): 1 year
- Cold storage (compliance): 7 years (regulatory requirement)

Best practices:

- Enable audit logging on all clusters (production mandatory)
- Use Metadata level minimum (RequestResponse for secrets)
- Separate control plane from application logs
- Implement log aggregation (don't rely on node storage)
- Encrypt logs at rest and in transit
- Implement log integrity protection (immutable storage)
- Regular log review and anomaly detection
- Automated alerting on security events

Understanding Kubernetes log types enables comprehensive security monitoring - control plane logs provide cluster-level visibility while node and application logs reveal runtime security events.

What are Kubernetes audit logs, how do you configure them, and what audit policy should you implement for security monitoring?

Kubernetes audit logs provide chronological record of all API server activities, essential for security monitoring, compliance, and forensic investigations.

Audit log fundamentals:

What audit logs capture:

- Every request to the kube-apiserver
- Who made the request (user, ServiceAccount, system component)
- What action was requested (verb: get, create, delete, etc.)
- Which resource was targeted
- When the request occurred
- Source IP and user agent
- Request payload (configurable)
- Response status and body (configurable)

Audit stages (lifecycle of request):

RequestReceived → ResponseStarted → ResponseComplete → Panic

- **RequestReceived:** Audit event generated as soon as request received, before any processing
- **ResponseStarted:** For long-running requests (watch), logged when headers sent but before body
- **ResponseComplete:** After response body sent
- **Panic:** Generated when request handler panicked

Audit levels (what to log):

- **None:** Don't log
- **Metadata:** Log request metadata (user, timestamp, resource, verb) but not request/response bodies
- **Request:** Log metadata and request body, not response
- **RequestResponse:** Log everything (metadata, request, and response bodies)

Configuring audit logging:

Method 1: kube-apiserver flags (most common):

```
# /etc/kubernetes/manifests/kube-apiserver.yaml (static pod)
apiVersion: v1
kind: Pod
metadata:
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
    - name: kube-apiserver
      image: registry.k8s.io/kube-apiserver:v1.28.0
      command:
        - kube-apiserver
        # Audit configuration
        - --audit-policy-file=/etc/kubernetes/audit-policy.yaml
        - --audit-log-path=/var/log/kubernetes/audit/audit.log
        - --audit-log-maxage=30          # Retain 30 days
        - --audit-log-maxbackup=10       # Keep 10 rotated files
        - --audit-log-maxsize=100        # Rotate at 100MB
        # Dynamic backend (webhook for real-time)
        - --audit-webhook-config-file=/etc/kubernetes/audit-webhook-config.yaml
        - --audit-webhook-initial-backoff=10s
  volumeMounts:
    - name: audit-policy
      mountPath: /etc/kubernetes/audit-policy.yaml
      readOnly: true
    - name: audit-log
      mountPath: /var/log/kubernetes/audit/
```

```

- name: audit-webhook
  mountPath: /etc/kubernetes/audit-webhook-config.yaml
  readOnly: true
volumes:
- name: audit-policy
  hostPath:
    path: /etc/kubernetes/audit-policy.yaml
    type: File
- name: audit-log
  hostPath:
    path: /var/log/kubernetes/audit/
    type: DirectoryOrCreate
- name: audit-webhook
  hostPath:
    path: /etc/kubernetes/audit-webhook-config.yaml
    type: File

```

Security-focused audit policy:

```

# /etc/kubernetes/audit-policy.yaml
apiVersion: audit.k8s.io/v1
kind: Policy
omitStages:
- RequestReceived # Don't log until request processed
rules:
# =====
# CRITICAL: Full logging for security resources
# =====

# Secrets - log everything (RequestResponse level)
- level: RequestResponse
  resources:
    - group: ""
      resources: ["secrets"]
  omitManagedFields: true # Reduce noise

# ServiceAccounts and tokens
- level: RequestResponse
  resources:
    - group: ""
      resources: ["serviceaccounts", "serviceaccounts/token"]

# RBAC resources (roles, bindings, etc.)
- level: RequestResponse
  resources:
    - group: "rbac.authorization.k8s.io"
      resources: ["roles", "rolebindings", "clusterroles", "clusterrolebindings"]

# Security-sensitive pod operations
- level: RequestResponse

```

```

verbs: ["create", "update", "patch"]
resources:
  - group: ""
    resources: ["pods", "pods/exec", "pods/portforward", "pods/proxy"]

# Network policies
- level: RequestResponse
  resources:
    - group: "networking.k8s.io"
      resources: ["networkpolicies"]

# Pod security policies (deprecated but still used)
- level: RequestResponse
  resources:
    - group: "policy"
      resources: ["podsecuritypolicies"]

# Admission webhooks (can bypass security)
- level: RequestResponse
  resources:
    - group: "admissionregistration.k8s.io"
      resources: ["mutatingwebhookconfigurations",
"validatingwebhookconfigurations"]

# =====
# HIGH: Metadata for important operations
# =====

# ConfigMaps (may contain sensitive data)
- level: Metadata
  resources:
    - group: ""
      resources: ["configmaps"]

# Persistent volumes (data access)
- level: Metadata
  resources:
    - group: ""
      resources: ["persistentvolumes", "persistentvolumeclaims"]

# Nodes (infrastructure)
- level: Metadata
  resources:
    - group: ""
      resources: ["nodes"]

# All creates, updates, deletes
- level: Metadata
  verbs: ["create", "update", "patch", "delete"]

# =====

```

```

# MEDIUM: Request level for write operations
# =====

# Deployments, StatefulSets, DaemonSets
- level: Request
  verbs: ["create", "update", "patch", "delete"]
  resources:
    - group: "apps"
      resources: ["deployments", "statefulsets", "daemonsets", "replicasets"]

# Jobs and CronJobs
- level: Request
  verbs: ["create", "update", "patch", "delete"]
  resources:
    - group: "batch"
      resources: ["jobs", "cronjobs"]

# =====
# LOW: Metadata for reads, None for noise
# =====

# Read-only operations on non-sensitive resources
- level: Metadata
  verbs: ["get", "list", "watch"]
  resources:
    - group: ""
      resources: ["endpoints", "services", "namespaces"]

# Exclude high-volume, low-value requests
- level: None
  users: ["system:kube-proxy"]
  verbs: ["watch"]
  resources:
    - group: ""
      resources: ["endpoints", "services"]

- level: None
  users: ["system:kube-controller-manager"]
  verbs: ["get", "update"]
  namespaces: ["kube-system"]

- level: None
  users: ["kubelet"]
  verbs: ["get"]
  resources:
    - group: ""
      resources: ["nodes"]

# System components health checks
- level: None
  userGroups: ["system:nodes"]

```

```

verbs: ["get"]
resources:
  - group: ""
    resources: ["nodes", "nodes/status"]

# Exclude read-only URLs
- level: None
nonResourceURLs:
  - /healthz*
  - /version
  - /swagger*

# =====
# DEFAULT: Catch-all for everything else
# =====

- level: Metadata

```

Webhook backend for real-time streaming:

```

# /etc/kubernetes/audit-webhook-config.yaml
apiVersion: v1
kind: Config
clusters:
- name: audit-webhook
  cluster:
    server: https://audit-collector.monitoring.svc.cluster.local:8443/api/v1/audit
    certificate-authority: /etc/kubernetes/pki/ca.crt
contexts:
- name: default
  context:
    cluster: audit-webhook
    user: audit-apiserver
current-context: default
users:
- name: audit-apiserver
  user:
    client-certificate: /etc/kubernetes/pki/audit-client.crt
    client-key: /etc/kubernetes/pki/audit-client.key

```

Managed Kubernetes audit configuration:

AKS (Azure Kubernetes Service):

```

# Enable diagnostic settings
az aks update \
--resource-group myResourceGroup \
--name myAKSCluster \
--enable-azure-rbac \

```

```
--enable-audit-logs

# Configure diagnostic settings to send to Log Analytics
az monitor diagnostic-settings create \
--resource /subscriptions/{sub-
id}/resourceGroups/myResourceGroup/providers/Microsoft.ContainerService/managedCluster-
s/myAKSCluster \
--name audit-logs-to-sentinel \
--workspace /subscriptions/{sub-
id}/resourceGroups/myResourceGroup/providers/Microsoft.OperationalInsights/workspaces/
sentinel-workspace \
--logs '['
{
  "category": "kube-audit",
  "enabled": true,
  "retentionPolicy": {
    "enabled": true,
    "days": 90
  }
},
{
  "category": "kube-audit-admin",
  "enabled": true,
  "retentionPolicy": {
    "enabled": true,
    "days": 90
  }
}
]'
]
```

EKS (Amazon Elastic Kubernetes Service):

```
# Enable control plane logging
aws eks update-cluster-config \
--name my-cluster \
--logging '{
  "clusterLogging": [
    {
      "types": ["api", "audit", "authenticator", "controllerManager", "scheduler"],
      "enabled": true
    }
  ]
}'

# Logs automatically sent to CloudWatch Logs
# Group: /aws/eks/my-cluster/cluster
```

GKE (Google Kubernetes Engine):

```
# Enable audit logging (enabled by default)
gcloud container clusters update my-cluster \
```

```
--enable-cloud-logging \
--logging=SYSTEM,WORKLOAD,API

# Logs automatically sent to Cloud Logging
# Can be exported to Cloud Storage, BigQuery, or Pub/Sub
```

Security monitoring queries (examples for analysis):

Detect secret access:

```
// Audit log query
{
  "objectRef.resource": "secrets",
  "verb": ["get", "list"],
  "user.username": {"$ne": "system:kube-controller-manager"}
}
```

Detect privilege escalation:

```
{
  "objectRef.resource": ["clusterrolebindings", "rolebindings"],
  "verb": ["create", "update", "patch"],
  "requestObject.roleRef.name": {"$in": ["cluster-admin", "admin"]}
}
```

Detect exec into pods:

```
{
  "objectRef.resource": "pods",
  "objectRef.subresource": "exec",
  "verb": "create",
  "responseStatus.code": 101 // Switching Protocols (successful)
}
```

Best practices:

- Start with conservative policy (Metadata level)
- Gradually increase to RequestResponse for critical resources
- Exclude high-volume, low-value events (health checks)
- Monitor audit log volume and adjust policy
- Use webhook backend for real-time SIEM integration
- Retain logs for compliance period (typically 1-7 years)
- Encrypt audit logs at rest
- Implement log integrity protection

- Regular review and tuning of audit policy
- Test policy changes in non-production first
- Document all policy decisions and exemptions

Proper audit configuration provides comprehensive security visibility while managing log volume and storage costs - essential for detecting threats, investigating incidents, and meeting compliance requirements.

How do you integrate Kubernetes logs (including control plane logs) with Microsoft Sentinel without using Grafana?

Integrating Kubernetes logs with Sentinel provides centralized security monitoring, correlation with other data sources, and advanced threat detection using Sentinel's analytics capabilities.

Integration architecture options:

Option 1: Azure Monitor Container Insights (AKS-native, recommended):

```
AKS Cluster → Container Insights Agent → Log Analytics Workspace → Sentinel
```

Option 2: Fluent Bit/Fluentd (flexible, multi-cloud):

```
K8s Cluster → Fluent Bit DaemonSet → Log Analytics → Sentinel
```

Option 3: Azure Event Hub (high-throughput):

```
K8s Cluster → Fluent Bit → Event Hub → Stream Analytics → Sentinel
```

Option 4: Direct API ingestion:

```
K8s Cluster → Custom Exporter → Data Collection API → Sentinel
```

Detailed implementation - Option 1: Azure Monitor Container Insights (AKS):

Enable Container Insights:

```
# Enable on existing AKS cluster
az aks enable-addons \
--resource-group myResourceGroup \
--name myAKSCluster \
--addons monitoring \
```

```
--workspace-resource-id /subscriptions/{sub-id}/resourceGroups/sentinel-rg/providers/Microsoft.OperationalInsights/workspaces/sentinel-workspace

# Verify agent deployment
kubectl get ds omsagent -n kube-system
kubectl get deployment omsagent-rs -n kube-system
```

Configure data collection:

```
# ConfigMap for Container Insights
apiVersion: v1
kind: ConfigMap
metadata:
  name: container-azm-ms-agentconfig
  namespace: kube-system
data:
  schema-version: v1
  config-version: ver1

# Log collection settings
log-data-collection-settings: |-
  [log_collection_settings]
    [log_collection_settings.stdout]
      enabled = true
      exclude_namespaces = ["kube-system", "kube-public"]

    [log_collection_settings.stderr]
      enabled = true
      exclude_namespaces = ["kube-system"]

    [log_collection_settings.env_var]
      enabled = true

    [log_collection_settings.enrich_container_logs]
      enabled = true

# Prometheus scraping (optional)
prometheus-data-collection-settings: |-
  [prometheus_data_collection_settings.cluster]
    interval = "1m"
    monitor_kubernetes_pods = true

  [prometheus_data_collection_settings.node]
    interval = "1m"
```

Enable control plane logs:

```
# Enable diagnostic settings for control plane
az monitor diagnostic-settings create \
```

```
--resource /subscriptions/{sub-
id}/resourceGroups/myResourceGroup/providers/Microsoft.ContainerService/managedClusters/myAKSCluster \
--name control-plane-logs \
--workspace /subscriptions/{sub-id}/resourceGroups/sentinel-
rg/providers/Microsoft.OperationalInsights/workspaces/sentinel-workspace \
--logs '['
{
  "category": "kube-apiserver",
  "enabled": true
},
{
  "category": "kube-audit",
  "enabled": true
},
{
  "category": "kube-audit-admin",
  "enabled": true
},
{
  "category": "kube-controller-manager",
  "enabled": true
},
{
  "category": "kube-scheduler",
  "enabled": true
},
{
  "category": "cluster-autoscaler",
  "enabled": true
},
{
  "category": "cloud-controller-manager",
  "enabled": true
},
{
  "category": "guard",
  "enabled": true
},
{
  "category": "csi-azuredisk-controller",
  "enabled": true
},
{
  "category": "csi-azurefile-controller",
  "enabled": true
}
]' \
--metrics '['
{
  "category": "AllMetrics",
```

```

        "enabled": true
    }
]'

```

Query control plane logs in Sentinel:

```

// kube-apiserver audit logs
AzureDiagnostics
| where Category == "kube-audit" or Category == "kube-audit-admin"
| where TimeGenerated > ago(24h)
| extend AuditLog = parse_json(log_s)
| extend
    User = tostring(AuditLog.user.username),
    Verb = tostring(AuditLog.verb),
    Resource = tostring(AuditLog.objectRef.resource),
    Namespace = tostring(AuditLog.objectRef.namespace),
    Name = tostring(AuditLog.objectRef.name),
    ResponseCode = toint(AuditLog.responseStatus.code),
    SourceIP = tostring(AuditLog.sourceIPs[0])
| project TimeGenerated, User, Verb, Resource, Namespace, Name, ResponseCode, SourceIP
| order by TimeGenerated desc

// kube-controller-manager logs
AzureDiagnostics
| where Category == "kube-controller-manager"
| where TimeGenerated > ago(1h)
| project TimeGenerated, log_s

// kube-scheduler logs
AzureDiagnostics
| where Category == "kube-scheduler"
| where TimeGenerated > ago(1h)
| project TimeGenerated, log_s

```

Detailed implementation - Option 2: Fluent Bit (works on any K8s):

Deploy Fluent Bit DaemonSet:

```

apiVersion: v1
kind: Namespace
metadata:
  name: logging
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: fluent-bit
  namespace: logging
---

```

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: fluent-bit-read
rules:
- apiGroups: []
  resources:
    - namespaces
    - pods
    - nodes
  verbs: ["get", "list", "watch"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: fluent-bit-read
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: fluent-bit-read
subjects:
- kind: ServiceAccount
  name: fluent-bit
  namespace: logging
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: fluent-bit-config
  namespace: logging
data:
  fluent-bit.conf: |
    [SERVICE]
    Flush      5
    Daemon     Off
    Log_Level  info
    Parsers_File parsers.conf

    # Tail container logs
    [INPUT]
    Name        tail
    Path        /var/log/containers/*.log
    Parser      docker
    Tag         kube.~
    Refresh_Interval 5
    Mem_Buf_Limit   50MB
    Skip_Long_Lines On

    # Tail control plane logs (if accessible)
    [INPUT]
    Name        tail

```

```

Path          /var/log/kube-apiserver-audit.log
Parser        json
Tag          audit.kube-apiserver
Mem_Buf_Limit 100MB

# Kubernetes metadata enrichment
[FILTER]
  Name      kubernetes
  Match     kube./*
  Kube_URL https://kubernetes.default.svc:443
  Kube_CA_File /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
  Kube_Token_File /var/run/secrets/kubernetes.io/serviceaccount/token
  Kube_Tag_Prefix kube.var.log.containers.
  Merge_Log On
  Keep_Log Off
  K8S-Logging.Parser On
  K8S-Logging.Exclude On

# Add cluster name
[FILTER]
  Name    modify
  Match   *
  Add    cluster_name production-cluster-01
  Add    environment production

# Output to Azure Log Analytics
[OUTPUT]
  Name      azure
  Match    *
  Customer_ID ${WORKSPACE_ID}
  Shared_Key ${WORKSPACE_KEY}
  Log_Type  KubernetesLogs
  Time_Generated true

parsers.conf: |
  [PARSER]
    Name      docker
    Format    json
    Time_Key  time
    Time_Format %Y-%m-%dT%H:%M:%S.%L%z

  [PARSER]
    Name      json
    Format    json
    Time_Key  timestamp
    Time_Format %Y-%m-%dT%H:%M:%S.%L%z
---

apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluent-bit

```

```

namespace: logging
labels:
  app: fluent-bit
spec:
  selector:
    matchLabels:
      app: fluent-bit
template:
  metadata:
    labels:
      app: fluent-bit
spec:
  serviceAccountName: fluent-bit
  containers:
    - name: fluent-bit
      image: fluent/fluent-bit:2.1.0
      env:
        - name: WORKSPACE_ID
          valueFrom:
            secretKeyRef:
              name: log-analytics-secret
              key: workspace-id
        - name: WORKSPACE_KEY
          valueFrom:
            secretKeyRef:
              name: log-analytics-secret
              key: workspace-key
  resources:
    limits:
      memory: 200Mi
    requests:
      cpu: 100m
      memory: 100Mi
  volumeMounts:
    - name: varlog
      mountPath: /var/log
      readOnly: true
    - name: varlibdockercontainers
      mountPath: /var/lib/docker/containers
      readOnly: true
    - name: fluent-bit-config
      mountPath: /fluent-bit/etc/
  volumes:
    - name: varlog
      hostPath:
        path: /var/log
    - name: varlibdockercontainers
      hostPath:
        path: /var/lib/docker/containers
    - name: fluent-bit-config
      configMap:

```

```
name: fluent-bit-config
```

Create Log Analytics secret:

```
# Get workspace credentials
WORKSPACE_ID=$(az monitor log-analytics workspace show \
--resource-group sentinel-rg \
--workspace-name sentinel-workspace \
--query customerId -o tsv)

WORKSPACE_KEY=$(az monitor log-analytics workspace get-shared-keys \
--resource-group sentinel-rg \
--workspace-name sentinel-workspace \
--query primarySharedKey -o tsv)

# Create Kubernetes secret
kubectl create secret generic log-analytics-secret \
--from-literal=workspace-id=$WORKSPACE_ID \
--from-literal=workspace-key=$WORKSPACE_KEY \
--namespace logging
```

Detailed implementation - Option 3: Control Plane Log Collection (self-managed clusters):

Collect audit logs with sidecar:

```
# If running kube-apiserver as static pod
apiVersion: v1
kind: Pod
metadata:
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
    # Main kube-apiserver container
    - name: kube-apiserver
      image: registry.k8s.io/kube-apiserver:v1.28.0
      command:
        - kube-apiserver
        - --audit-log-path=/var/log/kubernetes/audit.log
        - --audit-policy-file=/etc/kubernetes/audit-policy.yaml
      volumeMounts:
        - name: audit-log
          mountPath: /var/log/kubernetes

    # Sidecar to ship logs
    - name: audit-log-shipper
      image: fluent/fluent-bit:2.1.0
      env:
        - name: WORKSPACE_ID
```

```

valueFrom:
  secretKeyRef:
    name: log-analytics-secret
    key: workspace-id
- name: WORKSPACE_KEY
  valueFrom:
    secretKeyRef:
      name: log-analytics-secret
      key: workspace-key
volumeMounts:
- name: audit-log
  mountPath: /var/log/kubernetes
  readOnly: true
- name: fluent-bit-audit-config
  mountPath: /fluent-bit/etc/
volumes:
- name: audit-log
  hostPath:
    path: /var/log/kubernetes
    type: DirectoryOrCreate
- name: fluent-bit-audit-config
  configMap:
    name: fluent-bit-audit-config
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: fluent-bit-audit-config
  namespace: kube-system
data:
  fluent-bit.conf: |
    [SERVICE]
    Flush      5
    Log_Level  info

    [INPUT]
    Name        tail
    Path        /var/log/kubernetes/audit.log
    Parser      json
    Tag         audit.kube-apiserver
    Refresh_Interval 5
    Mem_Buf_Limit   100MB

    [FILTER]
    Name      modify
    Match    *
    Add      log_type kube-audit
    Add      cluster_name my-cluster

    [OUTPUT]

```

Name	azure
Match	*
Customer_ID	\${WORKSPACE_ID}
Shared_Key	\${WORKSPACE_KEY}
Log_Type	KubeAudit

Query Kubernetes logs in Sentinel:

```
// Container logs
KubernetesLogs_CL
| where TimeGenerated > ago(1h)
| where Namespace_s == "production"
| where ContainerName_s == "webapp"
| project TimeGenerated, Computer, PodName_s, ContainerName_s, LogEntry_s
| order by TimeGenerated desc

// Audit logs
KubeAudit_CL
| where TimeGenerated > ago(24h)
| extend Audit = parse_json(log_s)
| extend
    User = tostring(Audit.user.username),
    Verb = tostring(Audit.verb),
    Resource = tostring(Audit.objectRef.resource),
    ResponseCode = toint(Audit.responseStatus.code)
| where Resource == "secrets"
| project TimeGenerated, User, Verb, Resource, ResponseCode
```

Sentinel Analytics Rules (K8s-specific):

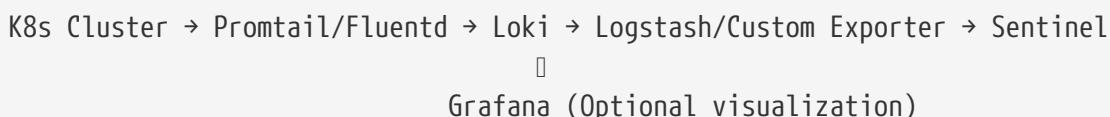
```
// Detect privileged pod creation
KubeAudit_CL
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| where tostring(Audit.verb) == "create"
| where tostring(Audit.objectRef.resource) == "pods"
| extend RequestObj = parse_json(tostring(Audit.requestObject))
| where RequestObj.spec.containers[0].securityContext.privileged == true
| project
    TimeGenerated,
    User = tostring(Audit.user.username),
    Namespace = tostring(Audit.objectRef.namespace),
    PodName = tostring(Audit.objectRef.name),
    SourceIP = tostring(Audit.sourceIPs[0])
```

Best practices for Kubernetes-Sentinel integration covered monitoring control plane, node, and application logs with proper enrichment, filtering, and security-focused analytics.

How do you integrate Kubernetes logs with Microsoft Sentinel using Grafana Loki as an intermediary, and what are the advantages of this approach?

Using Grafana Loki as a log aggregation layer before Sentinel provides additional benefits for log processing, cost optimization, and flexible querying while maintaining comprehensive security monitoring.

Architecture with Loki:



Why use Loki as intermediary:

Advantages:

- **Cost optimization** - filter and aggregate logs before sending to Sentinel (reduce ingestion costs), compress and deduplicate logs, selective forwarding of security-relevant events
- **Query flexibility** - LogQL for initial analysis and triage, Grafana dashboards for operational teams, KQL in Sentinel for security analytics
- **Buffer and resilience** - Loki acts as buffer during Sentinel outages, replay capability for missed events, local retention for fast queries
- **Multi-destination** - same logs to Sentinel (security) and other systems (operations), different retention policies per destination
- **Label-based filtering** - efficient filtering using Loki labels before expensive queries, reduce data sent to Sentinel

Detailed implementation:

Step 1: Deploy Loki in Kubernetes:

```
# Using Helm (recommended)  
# Add Grafana Helm repo  
helm repo add grafana https://grafana.github.io/helm-charts  
helm repo update  
  
# Create namespace  
kubectl create namespace monitoring  
  
# Install Loki with distributed mode for production
```

```

cat <<EOF > loki-values.yaml
loki:
  auth_enabled: false

  commonConfig:
    replication_factor: 3

  storage:
    type: 's3'
    bucketNames:
      chunks: loki-chunks
      ruler: loki-ruler
      admin: loki-admin
    s3:
      endpoint: s3.us-east-1.amazonaws.com
      region: us-east-1
      secretAccessKey: ${AWS_SECRET_KEY}
      accessKeyId: ${AWS_ACCESS_KEY}
      s3ForcePathStyle: false
      insecure: false

  schemaConfig:
    configs:
      - from: 2024-01-01
        store: tsdb
        object_store: s3
        schema: v12
        index:
          prefix: index_
          period: 24h

  limits_config:
    retention_period: 30d
    ingestion_rate_mb: 50
    ingestion_burst_size_mb: 100
    max_query_length: 30d
    max_query_parallelism: 32

# Resource allocation
read:
  replicas: 3
  resources:
    limits:
      cpu: 2
      memory: 4Gi
    requests:
      cpu: 1
      memory: 2Gi

write:
  replicas: 3
EOF

```

```

resources:
  limits:
    cpu: 2
    memory: 4Gi
  requests:
    cpu: 1
    memory: 2Gi

backend:
  replicas: 3
  resources:
    limits:
      cpu: 2
      memory: 4Gi
    requests:
      cpu: 1
      memory: 2Gi

gateway:
  enabled: true
  replicas: 2
  resources:
    limits:
      cpu: 1
      memory: 1Gi
    requests:
      cpu: 500m
      memory: 512Mi

# Security
serviceAccount:
  create: true
  annotations:
    eks.amazonaws.com/role-arn: arn:aws:iam::123456789012:role/loki-storage-role
EOF

helm install loki grafana/loki-distributed -f loki-values.yaml -n monitoring

```

Step 2: Deploy Promtail to ship logs to Loki:

```

cat <<EOF > promtail-values.yaml
config:
  clients:
    - url: http://loki-gateway.monitoring.svc.cluster.local/loki/api/v1/push

  snippets:
    scrapeConfigs: |
      # Scrape pod logs
      - job_name: kubernetes-pods
        pipeline_stages:

```

```

- cri: {}
- json:
  expressions:
    stream: stream
- labels:
  stream:
kubernetes_sd_configs:
- role: pod
relabel_configs:
# Only scrape pods with logging enabled
- source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_scrape]
  action: keep
  regex: true

# Add namespace label
- source_labels: [__meta_kubernetes_namespace]
  target_label: namespace

# Add pod name
- source_labels: [__meta_kubernetes_pod_name]
  target_label: pod

# Add container name
- source_labels: [__meta_kubernetes_pod_container_name]
  target_label: container

# Add node name
- source_labels: [__meta_kubernetes_pod_node_name]
  target_label: node

# Add security labels for filtering
- source_labels: [__meta_kubernetes_pod_label_security_monitoring]
  target_label: security_monitored

# Add environment
- source_labels: [__meta_kubernetes_namespace]
  target_label: environment
  regex: (.+)
  replacement: $1

# Scrape control plane logs (if accessible)
- job_name: kubernetes-control-plane
  static_configs:
- targets:
  - localhost
  labels:
    job: control-plane
    __path__: /var/log/kube-apiserver-audit.log
  pipeline_stages:
- json:
  expressions:

```

```

        user: user.username
        verb: verb
        resource: objectRef.resource
        namespace: objectRef.namespace
        responseCode: responseStatus.code
    - labels:
        user:
        verb:
        resource:
        namespace:
    - match:
        selector: '{job="control-plane"}'
        stages:
            - static_labels:
                log_type: audit
                security_critical: "true"

# DaemonSet for node coverage
daemonSet:
  enabled: true

# Security
serviceAccount:
  create: true

# Resources
resources:
  limits:
    cpu: 200m
    memory: 256Mi
  requests:
    cpu: 100m
    memory: 128Mi

# Volume mounts for log access
extraVolumes:
  - name: audit-logs
    hostPath:
      path: /var/log/kubernetes
      type: DirectoryOrCreate

extraVolumeMounts:
  - name: audit-logs
    mountPath: /var/log/kubernetes
    readOnly: true
EOF

helm install promtail grafana/promtail -f promtail-values.yaml -n monitoring

```

Step 3: Create Loki to Sentinel exporter:

```

# Custom Python service to query Loki and forward to Sentinel
import requests
import time
import json
import hashlib
from datetime import datetime, timedelta
from azure.monitor.ingestion import LogsIngestionClient
from azure.identity import DefaultAzureCredential

# Configuration
LOKI_URL = "http://loki-gateway.monitoring.svc.cluster.local:80"
SENTINEL_DCE = "https://dce-production.eastus-1.ingest.monitor.azure.com"
SENTINEL_DCR_ID = "dcr-xxxxxxxxxxxxxx"
SENTINEL_STREAM = "Custom-KubernetesLogs_CL"

# Initialize Sentinel client
credential = DefaultAzureCredential()
sentinel_client = LogsIngestionClient(
    endpoint=SENTINEL_DCE,
    credential=credential
)

# Track last processed timestamp
last_processed = {}

def query_loki(query, start_time, end_time):
    """Query Loki for logs"""
    params = {
        'query': query,
        'start': int(start_time.timestamp() * 1e9), # nanoseconds
        'end': int(end_time.timestamp() * 1e9),
        'limit': 5000
    }

    response = requests.get(
        f"{LOKI_URL}/loki/api/v1/query_range",
        params=params
    )

    if response.status_code == 200:
        return response.json()['data']['result']
    else:
        raise Exception(f"Loki query failed: {response.text}")

def transform_to_sentinel_format(loki_logs):
    """Transform Loki logs to Sentinel schema"""
    sentinel_logs = []

    for stream in loki_logs:
        labels = stream['stream']

```

```

        for entry in stream['values']:
            timestamp_ns, log_line = entry
            timestamp = datetime.fromtimestamp(int(timestamp_ns) / 1e9)

            # Parse JSON log if possible
            try:
                log_data = json.loads(log_line)
            except:
                log_data = {'message': log_line}

            sentinel_log = {
                'TimeGenerated': timestamp.isoformat() + 'Z',
                'Cluster': labels.get('cluster', 'unknown'),
                'Namespace': labels.get('namespace', ''),
                'Pod': labels.get('pod', ''),
                'Container': labels.get('container', ''),
                'Node': labels.get('node', ''),
                'LogLevel': log_data.get('level', 'INFO'),
                'Message': log_data.get('message', log_line),
                'Labels': json.dumps(labels),
                'LogData': json.dumps(log_data)
            }
            sentinel_logs.append(sentinel_log)

    return sentinel_logs

def forward_to_sentinel(logs, log_type):
    """Send logs to Sentinel"""
    if not logs:
        return

    try:
        sentinel_client.upload(
            rule_id=SENTINEL_DCR_ID,
            stream_name=SENTINEL_STREAM,
            logs=logs
        )
        print(f"Forwarded {len(logs)} {log_type} logs to Sentinel")
    except Exception as e:
        print(f"Failed to forward logs: {str(e)}")

def process_security_logs():
    """Process security-critical logs from Loki"""
    end_time = datetime.utcnow()
    start_time = last_processed.get('security', end_time - timedelta(minutes=5))

    # Query 1: Audit logs
    audit_query = '{log_type="audit", security_critical="true"}'
    audit_logs = query_loki(audit_query, start_time, end_time)

```

```

audit_sentinel = transform_to_sentinel_format(audit_logs)
forward_to_sentinel(audit_sentinel, 'audit')

# Query 2: Security-monitored pods
security_query = '{security_monitored="true"} |= "error" or "failed" or
"unauthorized" or "forbidden"'
security_logs = query_loki(security_query, start_time, end_time)
security_sentinel = transform_to_sentinel_format(security_logs)
forward_to_sentinel(security_sentinel, 'security')

# Query 3: Privileged container logs
privileged_query = '{namespace=~"production|staging"} | json | privileged="true"'
privileged_logs = query_loki(privileged_query, start_time, end_time)
privileged_sentinel = transform_to_sentinel_format(privileged_logs)
forward_to_sentinel(privileged_sentinel, 'privileged')

last_processed['security'] = end_time

def process_error_logs():
    """Process error logs from Loki"""
    end_time = datetime.utcnow()
    start_time = last_processed.get('errors', end_time - timedelta(minutes=5))

    # Query for errors across all pods
    error_query = '{namespace!~"kube-system|kube-public"} |~
    "(?i)error|exception|fatal"'
    error_logs = query_loki(error_query, start_time, end_time)
    error_sentinel = transform_to_sentinel_format(error_logs)
    forward_to_sentinel(error_sentinel, 'errors')

    last_processed['errors'] = end_time

def main():
    """Main processing loop"""
    print("Starting Loki to Sentinel forwarder...")

    while True:
        try:
            # Process different log types
            process_security_logs()
            process_error_logs()

            # Sleep before next iteration
            time.sleep(60) # Run every minute

        except Exception as e:
            print(f"Error in processing loop: {str(e)}")
            time.sleep(60)

if __name__ == "__main__":

```

```
main()
```

Deploy exporter as Kubernetes Deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: loki-sentinel-exporter
  namespace: monitoring
spec:
  replicas: 2
  selector:
    matchLabels:
      app: loki-sentinel-exporter
  template:
    metadata:
      labels:
        app: loki-sentinel-exporter
    spec:
      serviceAccountName: loki-sentinel-exporter
      containers:
        - name: exporter
          image: myregistry.azurecr.io/loki-sentinel-exporter:latest
          env:
            - name: LOKI_URL
              value: "http://loki-gateway.monitoring.svc.cluster.local:80"
            - name: AZURE_CLIENT_ID
              valueFrom:
                secretKeyRef:
                  name: sentinel-credentials
                  key: client-id
            - name: AZURE_TENANT_ID
              valueFrom:
                secretKeyRef:
                  name: sentinel-credentials
                  key: tenant-id
            - name: AZURE_CLIENT_SECRET
              valueFrom:
                secretKeyRef:
                  name: sentinel-credentials
                  key: client-secret
      resources:
        limits:
          cpu: 500m
          memory: 512Mi
        requests:
          cpu: 200m
          memory: 256Mi
---
apiVersion: v1
```

```

kind: ServiceAccount
metadata:
  name: loki-sentinel-exporter
  namespace: monitoring
  annotations:
    azure.workload.identity/client-id: "your-managed-identity-client-id"

```

Step 4: Query and analyze in Sentinel:

```

// Query forwarded Kubernetes logs
KubernetesLogs_CL
| where TimeGenerated > ago(1h)
| extend Labels = parse_json(Labels_s)
| extend LogData = parse_json(LogData_s)
| where Namespace_s == "production"
| project
  TimeGenerated,
  Cluster_s,
  Namespace_s,
  Pod_s,
  Container_s,
  LogLevel_s,
  Message_s,
  Labels,
  LogData
| order by TimeGenerated desc

// Detect suspicious activity
KubernetesLogs_CL
| where TimeGenerated > ago(5m)
| where Message_s contains "error" or Message_s contains "failed"
| extend Labels = parse_json(Labels_s)
| where Labels.security_monitored == "true"
| summarize
  ErrorCount = count(),
  Pods = make_set(Pod_s)
  by Namespace_s, bin(TimeGenerated, 5m)
| where ErrorCount > 10

```

Advanced filtering with Loki:

```

# Loki configuration for selective forwarding
# Only forward security-relevant logs to reduce Sentinel costs

limits_config:
  # Retention in Loki (cheap storage)
  retention_period: 30d

  # Rate limits

```

```

ingestion_rate_mb: 50
ingestion_burst_size_mb: 100

# Compactor for cost optimization
compactor:
  working_directory: /data/compactor
  shared_store: s3
  compaction_interval: 10m
  retention_enabled: true
  retention_delete_delay: 2h
  retention_delete_worker_count: 150

# Query for selective export
# LogQL query in exporter:
# {namespace=~"production|staging", security_monitored="true"}
# |= "error" or "unauthorized" or "exec" or "secret"
# Only these logs sent to Sentinel

```

Cost comparison:

Without Loki (direct to Sentinel):

- All logs: 500 GB/day
- Sentinel ingestion: $500 \text{ GB} \times \$2.30/\text{GB} = \$1,150/\text{day} = \$34,500/\text{month}$

With Loki (filtered):

- All logs to Loki: 500 GB/day
- Loki storage (S3): $500 \text{ GB} \times \$0.023/\text{GB} = \$11.50/\text{day} = \$345/\text{month}$
- Filtered to Sentinel: 50 GB/day (10% security-critical)
- Sentinel ingestion: $50 \text{ GB} \times \$2.30/\text{GB} = \$115/\text{day} = \$3,450/\text{month}$
- Total: \$3,795/month

Savings: \$30,705/month (89% reduction)

Benefits summary:

- **Cost optimization** - 80-90% reduction in Sentinel ingestion costs
- **Query performance** - Fast queries in Loki for operational issues
- **Flexibility** - LogQL for developers, KQL for security team
- **Resilience** - Loki buffers during Sentinel outages
- **Multi-tenant** - Different teams query Loki, security uses Sentinel

Best practices:

- Use Loki labels efficiently (avoid high cardinality)
- Implement retention policies (7 days hot, 30 days warm in Loki)
- Filter aggressively before Sentinel (security-critical only)

- Monitor exporter performance and lag
- Implement alerting for export failures
- Regular review of forwarding rules
- Test query performance in both systems
- Document which logs go where

Using Loki as intermediary provides cost-effective log management while maintaining comprehensive security monitoring in Sentinel - ideal for large-scale Kubernetes deployments where full log ingestion to SIEM would be prohibitively expensive.

What are the key security events you should detect and alert on from Kubernetes control plane logs in Sentinel?

Effective security monitoring requires detection rules targeting specific attack patterns and policy violations visible in control plane logs.

Critical detection scenarios:

1. Unauthorized API access attempts:

```
// Detect failed authentication attempts
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| extend
  User = tostring(Audit.user.username),
  ResponseCode = toint(Audit.responseStatus.code),
  Reason = tostring(Audit.responseStatus.reason),
  SourceIP = tostring(Audit.sourceIPs[0])
| where ResponseCode in (401, 403) // Unauthorized or Forbidden
| summarize
  FailedAttempts = count(),
  Resources = make_set(tostring(Audit.objectRef.resource)),
  FirstAttempt = min(TimeGenerated),
  LastAttempt = max(TimeGenerated)
  by User, SourceIP, Reason
| where FailedAttempts > 5
| extend Severity = case(
  FailedAttempts > 20, "High",
  FailedAttempts > 10, "Medium",
  "Low"
)
```

2. Secret access monitoring:

```
// Alert on secret enumeration or mass access
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(10m)
| extend Audit = parse_json(log_s)
| where tostring(Audit.objectRef.resource) == "secrets"
| where tostring(Audit.verb) in ("get", "list")
| extend
    User = tostring(Audit.user.username),
    Namespace = tostring(Audit.objectRef.namespace),
    SecretName = tostring(Audit.objectRef.name),
    SourceIP = tostring(Audit.sourceIPs[0])
| summarize
    UniqueSecrets = dcount(SecretName),
    Namespaces = make_set(Namespace),
    Secrets = make_set(SecretName)
    by User, SourceIP, bin(TimeGenerated, 5m)
| where UniqueSecrets > 10 // Accessed >10 secrets in 5 min
| extend
    AlertTitle = strcat("Secret Enumeration: ", User, " accessed ", UniqueSecrets, " secrets"),
    Severity = "High"
```

3. Privilege escalation detection:

```
// Detect creation of privileged roles or bindings
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| where tostring(Audit.objectRef.resource) in ("clusterrolebindings", "rolebindings")
| where tostring(Audit.verb) in ("create", "update", "patch")
| extend
    User = tostring(Audit.user.username),
    BindingName = tostring(Audit.objectRef.name),
    RequestObj = parse_json(tostring(Audit.requestObject))
| extend RoleName = tostring(RequestObj.roleRef.name)
| where RoleName in ("cluster-admin", "admin", "edit") // Privileged roles
| project
    TimeGenerated,
    User,
    BindingName,
    RoleName,
    Namespace = tostring(Audit.objectRef.namespace),
    Subjects = RequestObj.subjects,
    SourceIP = tostring(Audit.sourceIPs[0])
| extend
```

```

AlertTitle = strcat("Privilege Escalation: ", User, " granted ", RoleName, " to ",
toString(Subjects[0].name)),
Severity = "Critical"

```

4. Exec/Port-forward sessions:

```

// Monitor interactive access to containers
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| where toString(Audit.objectRef.subresource) in ("exec", "portforward", "attach")
| where toString(Audit.verb) == "create"
| where toint(Audit.responseStatus.code) == 101 // Switching Protocols (successful)
| extend
    User = toString(Audit.user.username),
    PodName = toString(Audit.objectRef.name),
    Namespace = toString(Audit.objectRef.namespace),
    Subresource = toString(Audit.objectRef.subresource),
    SourceIP = toString(Audit.sourceIPs[0])
| project
    TimeGenerated,
    User,
    Namespace,
    PodName,
    Subresource,
    SourceIP
| extend
    AlertTitle = strcat("Container Access: ", User, " executed ", Subresource, " on ",
Namespace, "/", PodName),
    Severity = case(
        Namespace in ("kube-system", "production"), "High",
        "Medium"
    )

```

5. Privileged container creation:

```

// Detect creation of containers with elevated privileges
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| where toString(Audit.objectRef.resource) == "pods"
| where toString(Audit.verb) == "create"
| extend RequestObj = parse_json(tostring(Audit.requestObject))
| mv-expand Container = RequestObj.spec.containers
| extend SecurityContext = Container.securityContext
| where
    SecurityContext.privileged == true

```

```

or SecurityContext.capabilities.add has "SYS_ADMIN"
or SecurityContext.capabilities.add has "NET_ADMIN"
or Container.image startswith "/" // hostPath mount
| project
    TimeGenerated,
    User = tostring(Audit.user.username),
    Namespace = tostring(Audit.objectRef.namespace),
    PodName = tostring(Audit.objectRef.name),
    ContainerName = tostring(Container.name),
    Image = tostring(Container.image),
    Privileged = SecurityContext.privileged,
    Capabilities = SecurityContext.capabilities,
    SourceIP = tostring(Audit.sourceIPs[0])
| extend
    AlertTitle = strcat("Privileged Container: ", User, " created privileged pod ",
Namespace, "/", PodName),
    Severity = "High"

```

6. Admission controller bypasses:

```

// Detect modifications to admission controllers
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| where tostring(Audit.objectRef.resource) in ("mutatingwebhookconfigurations",
"validatingwebhookconfigurations")
| where tostring(Audit.verb) in ("update", "patch", "delete")
| extend
    User = tostring(Audit.user.username),
    WebhookName = tostring(Audit.objectRef.name),
    Action = tostring(Audit.verb)
| project
    TimeGenerated,
    User,
    WebhookName,
    Action,
    SourceIP = tostring(Audit.sourceIPs[0])
| extend
    AlertTitle = strcat("Security Control Modified: ", User, " ", Action, " admission
webhook ", WebhookName),
    Severity = "Critical"

```

7. Suspicious API usage patterns:

```

// Detect unusual API call patterns (reconnaissance)
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(10m)

```

```

| extend Audit = parse_json(log_s)
| extend
    User = tostring(Audit.user.username),
    Verb = tostring(Audit.verb),
    Resource = tostring(Audit.objectRef.resource)
| where Verb in ("list", "get") // Read operations
| summarize
    UniqueResources = dcount(Resource),
    Resources = make_set(Resource),
    RequestCount = count()
    by User, bin(TimeGenerated, 5m)
| where UniqueResources > 15 // Accessed >15 different resource types
| extend
    AlertTitle = strcat("API Reconnaissance: ", User, " queried ", UniqueResources, " resource types"),
    Severity = "Medium"

```

8. After-hours activity:

```

// Detect administrative actions outside business hours
let BusinessHours = dynamic(["09", "10", "11", "12", "13", "14", "15", "16", "17"]);
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(1h)
| extend Audit = parse_json(log_s)
| extend Hour = format_datetime(TimeGenerated, "HH")
| where Hour !in (BusinessHours)
| where tostring(Audit.verb) in ("create", "update", "patch", "delete")
| where tostring(Audit.objectRef.resource) in ("secrets", "rolebindings",
"clusterrolebindings", "pods")
| extend
    User = tostring(Audit.user.username),
    Resource = tostring(Audit.objectRef.resource),
    Action = tostring(Audit.verb),
    SourceIP = tostring(Audit.sourceIPs[0])
| where User !startswith "system:" // Exclude system accounts
| project
    TimeGenerated,
    Hour,
    User,
    Resource,
    Action,
    SourceIP
| extend
    AlertTitle = strcat("After-Hours Activity: ", User, " ", Action, " ", Resource, " at ",
Hour, ":00"),
    Severity = "Medium"

```

9. Anonymous or unauthenticated access:

```

// Detect anonymous API access attempts
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| extend User = tostring(Audit.user.username)
| where User in ("system:anonymous", "system:unauthenticated")
| where toint(Audit.responseStatus.code) == 200 // Successful
| extend
    Resource = tostring(Audit.objectRef.resource),
    Verb = tostring(Audit.verb),
    SourceIP = tostring(Audit.source IPs[0])
| project
    TimeGenerated,
    User,
    Resource,
    Verb,
    SourceIP
| extend
    AlertTitle = "Anonymous API Access Allowed",
    Severity = "High"

```

10. ServiceAccount token theft detection:

```

// Detect ServiceAccount tokens used from unexpected locations
let KnownPodIPs =
    // Build table of known pod IPs (from Container Insights)
    KubePodInventory
    | where TimeGenerated > ago(1h)
    | summarize by PodIp, ServiceName
);
AzureDiagnostics
| where Category == "kube-audit"
| where TimeGenerated > ago(5m)
| extend Audit = parse_json(log_s)
| extend User = tostring(Audit.user.username)
| where User startswith "system:serviceaccount:"
| extend SourceIP = tostring(Audit.source IPs[0])
| join kind=leftanti (KnownPodIPs) on $left.SourceIP == $right.PodIp
| where SourceIP !startswith "10." // Not from cluster network
| project
    TimeGenerated,
    ServiceAccount = User,
    SourceIP,
    Resource = tostring(Audit.objectRef.resource),
    Verb = tostring(Audit.verb)
| extend
    AlertTitle = strcat("ServiceAccount Token Misuse: ", ServiceAccount, " used from
external IP ", SourceIP),

```

```
Severity = "Critical"
```

Best practices for detection rules:

- Tune thresholds based on baseline (reduce false positives)
- Include context in alerts (user, IP, resource, action)
- Severity based on risk (secrets > configmaps)
- Exclude expected system activity (kube-controller-manager, kubelet)
- Alert aggregation (don't spam on every event)
- Include remediation guidance in alerts
- Regular review and tuning of rules
- Test rules against historical data
- Document rule logic and expected false positive rate

These detection rules transform raw audit logs into actionable security intelligence, enabling rapid threat detection and response in Kubernetes environments.